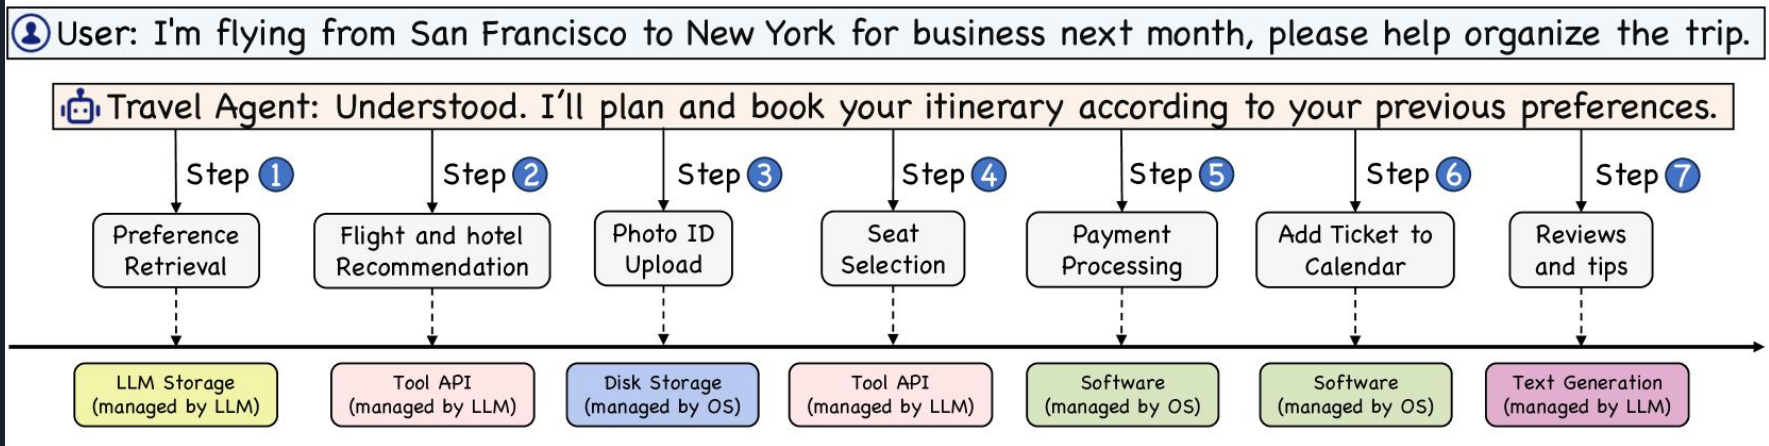
A decorative graphic on the left side of the slide consisting of two overlapping parallelograms. The front one is blue and the back one is a light green. They are positioned diagonally, with the blue one partially covering the green one.

AIOS: LLM Agent Operating System Short Story Review

By Shawn Chumbar

Introduction

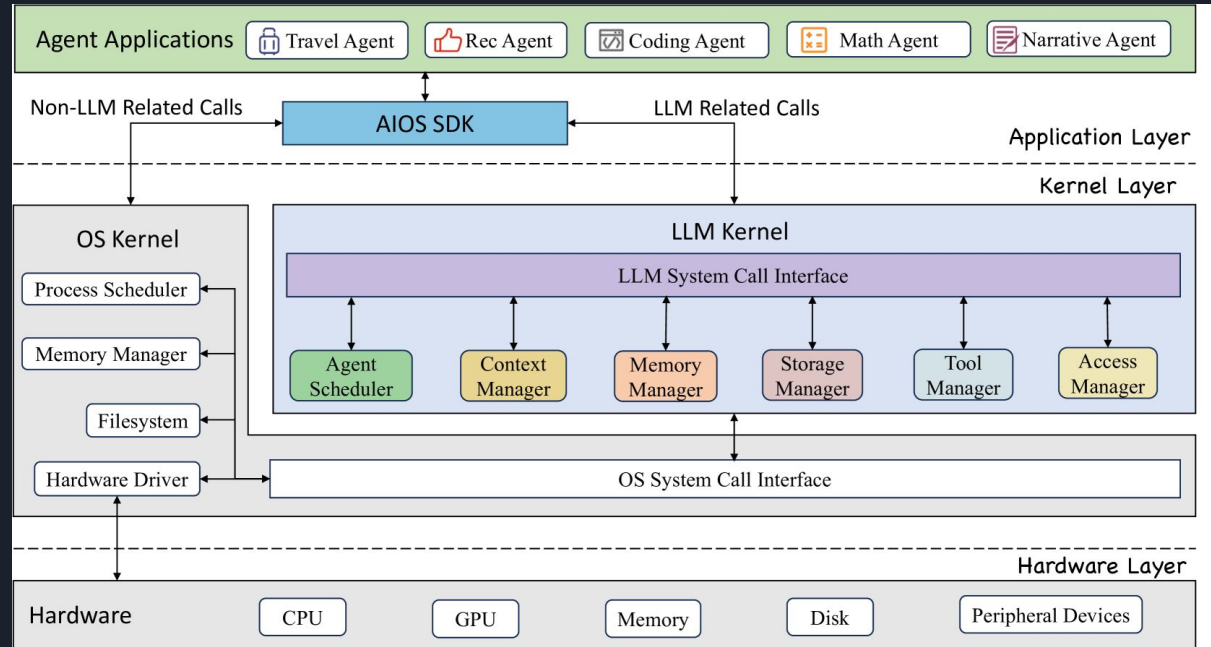
- What is AIOS?
- Objective of AIOS
- Challenges addressed by AIOS



AIOS Architecture Overview

Three Distinct Layers:

1. Application
2. Hardware
3. Kernel





Application Layer

- Agent development using AIOS SDK
- Abstracting complexities of lower-level system functions



Hardware Layer

- Physical components:
 - CPU, GPU, Memory, Disk, etc.
- LLM Kernel leveraging Hardware resources without direct management

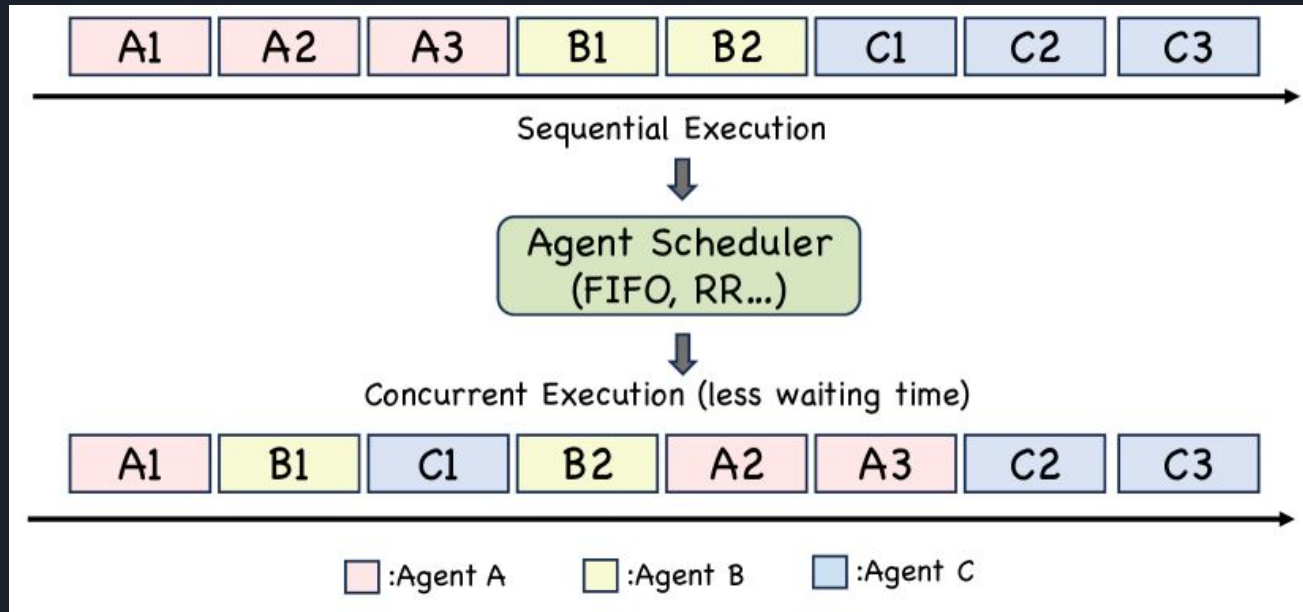


Kernel Layer

- OS Kernel for non-LLM tasks
- LLM kernel for LLM-specific operations
- This includes the following modules:
 1. Agent Scheduler
 2. Context Manager
 3. Memory Manager
 4. Storage Manager
 5. Tool Manager
 6. Access Manager

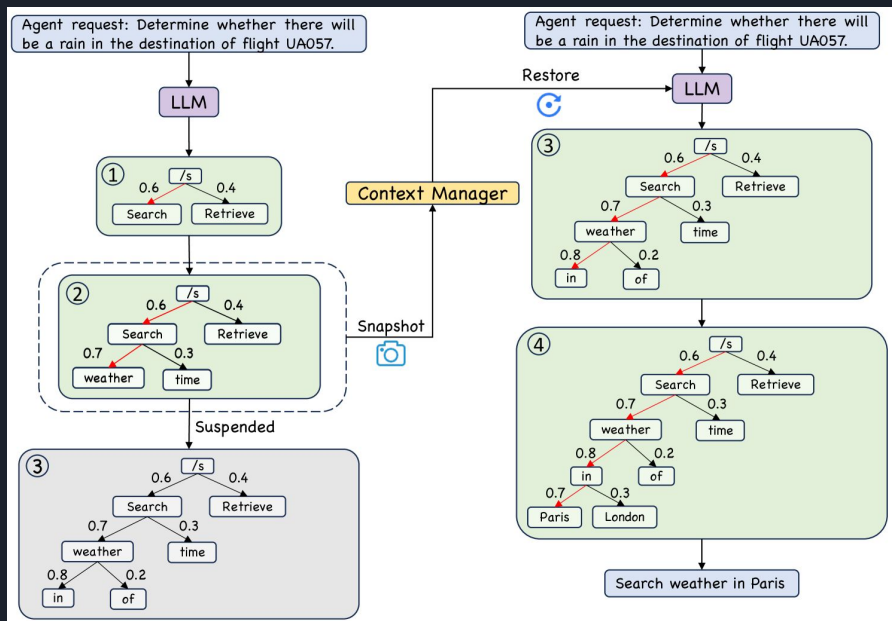
Agent Scheduler

- Efficient management of agent requests
- Scheduling algorithms
 - FIFO, Round Robin, etc.
- Balancing waiting time and turnaround time for parallel agent execution



Context Manager

- Critical functions
 - Context snapshot and restoration
 - Context window management
- Example of context snapshot usage and restoration process
- Context window management includes text summarization and expansion techniques for context window management





Memory Manager

- Memory Manager
 - Short-term memory management within an agent's lifecycle
 - Independent storage of each agent's memory



Storage Manager

- Storage Manager
 - Long-term data preservation (local files, databases, cloud-based storage)
 - Retrieval augmentation and enriching agent knowledge update



Tool Manager

- Management of diverse API tools for LLM functionality enhancement
- Examples of API calls (web search, scientific computing, database retrieval, image processing)



Access Manager

- Access control operations among agents
- Privilege groups, audit logs, and protection against potential privilege attacks



LLM System Call Interface and AIOS SDK

- Basic LLM call operation functions bridging complex agent requests and kernel module execution
- AIOS SDK: Toolkit for creating agent applications within AIOS
- Functionalities for agent initialization, lifecycle management, resource management, and task planning



Future Work

- Advanced scheduling algorithms for dependency analysis and resource optimization
- Improved context management efficiency (time and space-efficient solutions)
- Enhanced safety and privacy measures against malicious attacks and for data safeguarding



Personal Analysis

- Significance of AIOS in intelligent agent management and execution
- Standout features: Modular architecture, LLM kernel segregation, agent scheduler, context manager
- Future potential and impact on AI-powered computing



Conclusion

- AIOS as a milestone in the evolution of operating systems and AI agent integration
- Potential for revolutionizing intelligent agent development, deployment, and interaction
- Anticipation of future advancements and challenges in AI operating systems



Works Cited

Mei, K., Li, Z., Xu, S., Ye, R., Ge, Y., & Zhang, Y. (2024). AIOS: LLM Agent Operating System. Retrieved from arXiv preprint server (ID: 2403.16971).

WorldofAI. "AIOS: LLM Agent Operating System! Create Software, Automate Tasks, and More!" YouTube, 10 Apr. 2024,
https://www.youtube.com/watch?v=7WVP0yWVrLI&ab_channel=WorldofAI.